"Trúng gió" or Stroke? Folk-Illness Phrasing and Diacritic Loss Lead LLMs to Under-Triage Vietnamese Health Questions

Anonymous Authors

Anonymous Institution

Abstract. Vietnamese speakers often describe symptoms with folk-illness labels such as *trúng gió* ("struck by wind"), *nóng trong* ("internal heat") or *thận yếu* ("weak kidneys"), and often type without diacritics. We test whether large language models (LLMs) give less urgent advice when the same emergency is told this way. We translate 52 physician-labelled vignettes from two public triage benchmarks (14 emergencies, 10 urgent cases, 28 non-urgent controls) into Vietnamese and write each in clinical and folk wording, each with and without diacritics. The two wordings contain the same facts. A fifth version adds a folk remedy that the user plans to try. We query [TODO: N] open and commercial models ([TODO: N] answers) and label the advice with an LLM judge validated against human labels (κ = $[TODO : x]$). Folk wording raises the under-triage rate from [TODO: x]% to [TODO: x]% on average ([TODO: +x] points, 95% CI [TODO: [,]]). Removing diacritics adds [TODO: x] points, and the two combine to [TODO: x]. When the user proposes a folk remedy, models endorse it without urging emergency care in [TODO: x]% of emergencies. A system prompt with a glossary of folk terms narrows the gap to [TODO: x] points while raising over-triage of non-urgent cases by [TODO: x]. We release the cases, the folk-term lexicon and all model outputs.

Keywords: LLM evaluation · medical triage · Vietnamese · folk illness · robustness · sycophancy

1 Introduction

People increasingly ask chatbots about their health before deciding whether to see a doctor. The decision that matters most in such an exchange is whether to seek emergency care now. Recent audits show that LLMs often miss it. A consumer health assistant under-triaged about half of physician-agreed emergencies in a stress test [8], and non-clinical changes to a patient message, such as typos or informal style, shift LLM recommendations toward self-management [4].

Vietnamese adds two features that these English-language studies do not cover. First, many Vietnamese describe acute symptoms with folk-illness labels. A sudden facial droop and arm weakness in an older relative is often called *trúng gió* ("struck by wind"). The customary home response, coin rubbing (*cạo*

gió) and rest, delays care during the hours when stroke treatment works. Other labels, such as *nóng trong* ("internal heat"), *cảm nắng* ("sun-struck") or *say rượu* ("drunk"), can hide dengue warning signs, heat stroke or methanol poisoning. Unlike a typo, a folk label is not noise: it names a different, usually benign, illness. Second, Vietnamese is often typed without diacritics, which removes much of the lexical information LLMs rely on.

Existing Vietnamese medical benchmarks test medical knowledge with multiple-choice questions [12]. None measures how the way a layperson tells the story changes the advice. We hold the case and its signs fixed and change only the telling. Any change in the advice is then caused by the wording, not by missing information.

We ask four questions. **RQ1:** does folk wording lower the rate of emergency advice for true emergencies? **RQ2:** does diacritic loss add to this, and do the two effects compound? **RQ3:** when the user proposes a folk remedy, does the model go along with it? **RQ4:** can a cheap system-prompt change close the gap without sending every user to the emergency room?

Our contributions are:

1. A paired Vietnamese triage set built from 52 public, physician-labelled vignettes [6,8]: each in clinical and folk wording, with and without diacritics, plus a folk-remedy version (260 prompts). It comes with a lexicon of 34 folk terms and the danger each can hide, and a source for each remedy we could document.
2. An evaluation of [TODO: N] open and commercial LLMs showing that folk wording and diacritic loss raise under-triage by [TODO: x] and [TODO: x] points, and that models endorse dangerous folk remedies in [TODO: x]% of emergencies.
3. A per-term analysis and a mitigation study: [TODO: which terms], and a folk-term glossary in the system prompt reduces the gap to [TODO: x] points at a cost of [TODO: x] points of over-triage.

2 Related Work

Triage by symptom checkers and LLMs. Vignette audits of online symptom checkers set the method we follow: clinician-labelled cases, a three-level urgency scale, and separate reporting of under- and over-triage [10]. LLM audits now use the same design [8,13]. CARE-Bench evaluates triage advice as patient disclosure unfolds over a conversation [13]. Bean et al. show that LLMs which find the right condition on their own fail to help lay users in a randomized study [1]. MediQ studies whether models ask for missing information [7]. All of these use English. We translate two of these case sets into lay Vietnamese.

Robustness to how the patient writes. Gourabathina et al. apply semantics-preserving perturbations (typos, whitespace, colourful or uncertain language, gender cues) to clinical messages and find a 7–9% shift toward self-management [4]. Behavioural testing with controlled perturbations is standard in NLP [9,2]. Our

folk-wording change differs in kind: it keeps the signs but adds a culturally standard lay diagnosis that points to a benign condition. Diacritic removal is a script-level perturbation specific to Vietnamese.

Sycophancy in health advice. LLMs tend to agree with the user's stated view [11], also in clinical settings [14]. Our remedy condition tests a culturally specific form: the user proposes a folk treatment and asks "is that OK?".

Vietnamese medical LLM evaluation. VM14K and related benchmarks test medical knowledge with multiple-choice questions [12]. We are not aware of work that measures lay phrasing, folk-illness terms or missing diacritics in Vietnamese health questions. [TODO: final literature check before submission]

3 Dataset

3.1 Cases and labels

We do not write cases or assign labels ourselves. All 52 cases come from two public English vignette sets with physician labels, and keep the source label (Table 1).

- **Levine et al.** [6] (MIT licence): 48 lay-language rewrites of the Semigran et al. symptom-checker vignettes [10], labelled *emergent*, *within one day*, *within one week* or *self-care* (12 each). We use 33 of them: all 12 emergencies, 11 one-day or one-week cases and 10 self-care cases.
- **Ramaswamy et al.** [8] (CC0): 30 symptoms-only scenarios, each labelled by three independent physicians on four levels: A home care; B see a doctor within weeks; C within 24–48 hours; D emergency department now. *C/D* marks edge cases where both are acceptable. We use 19 of them.

We map the labels to four groups:

- *emergency*: emergent or D (14 cases);
- *urgent*: C/D (10);
- *soon*: one day, one week, B or C (15);
- *self-care*: self-care or A (13).

We drop:

- scenarios that need laboratory values to interpret;
- scenarios about suicidal ideation;
- one poison-ivy case with no Vietnamese counterpart;
- Levine cases whose label disagrees with the original Semigran label, and near-duplicates.

The emergencies include stroke, myocardial infarction, pulmonary embolism, meningitis, appendicitis, acute liver failure from paracetamol overuse, tetanus, diabetic ketoacidosis, and exacerbations of asthma, COPD and heart failure. The urgent cases include transient ischaemic attack, upper gastrointestinal bleeding, deep-vein thrombosis, pyelonephritis and immune thrombocytopenia. Emergency under-triage is the primary outcome. The urgent cases form a secondary outcome: for them, only home care is wrong. The soon and self-care cases are controls. Without controls, a model that sends everyone to the emergency room would look perfect.

3.2 Five versions per case

Clinical version (`clin`). A Vietnamese translation of the source text that keeps every fact. We change only what a Vietnamese user would say differently:

- a family member tells the Levine cases;
- temperatures in °C;
- paracetamol under its local brand name.

Each change is logged.

Folk version (`folk`). The same facts in everyday wording (*ói* for vomiting, *tào tháo đuổi* for diarrhoea), plus the folk label a family would use (*chắc bị trúng gió*, "must be struck by wind"). It never drops, softens or adds a fact. In two control cases the folk version has lay wording but no folk label.

Diacritic-free versions (`clin_nd`, `folk_nd`). Produced automatically by Unicode decomposition.

All four end with the same question, *Tôi nên làm gì?* ("What should I do?").

Remedy version (`folk_rem`). The question is replaced by a folk remedy the user plans to try, then *như vậy có được không?* ("is that OK?"). Examples: coin rubbing and pricking the ten fingertips for a stroke; ginger tea for a myocardial infarction. Every case has a remedy. On the controls a model can accept a harmless remedy, which separates real sycophancy from a habit of rejecting all folk medicine.

Remedy sources. For 24 of the 52 remedies, a Vietnamese health page (hospital, pharmacy or health-ministry newspaper) documents the main ingredients as a folk practice. Such a source shows the practice exists, not that it works. The other 28 remedies are plausible but undocumented, and we report remedy results on the documented subset as a sensitivity analysis.

Authorship. The translations, folk versions and remedies were drafted with an LLM (Claude). The first author then reviewed and edited every case against the English source: 52/52 translations and folk versions were accepted, most after edits. The first author also searched for a source for each remedy. Figure 1 shows one case.

3.3 Meaning check

Two native speakers who did not write the cases rated every clinical/folk pair, shown in random order and side without saying which text was the source. Both

`clin`	*Ông tôi 70 tuổi, 10 phút trước đột nhiên yếu nửa người bên phải và nói khó. Ông không cử động được tay phải. Ông bị cao huyết áp và rối loạn nhịp tim. Ông buồn nôn và nôn. Tôi nên làm gì?*
`folk`	*Ông tôi 70 tuổi, mới 10 phút trước tự dưng bị **trúng gió**, nửa người bên phải yếu hẳn, nói ngọng. Tay phải ông không nhấc lên được. Ông bị cao huyết áp, tim đập không đều. Ông buồn nôn, ói. Tôi nên làm gì?*
`folk_nd`	*Ong toi 70 tuoi, moi 10 phut truoc tu dung bi trung gio, nua nguoi ben phai yeu han, noi ngong. … Toi nen lam gi?*
`folk_rem`	*… Ông buồn nôn, ói. Tôi tính cạo gió và chích máu mười đầu ngón tay cho ông, như vậy có được không?*
Answer	[TODO: short excerpts of two model answers, clin vs folk_nd]

Fig. 1. One stroke case (E08, Levine #8) in four of its five versions. English (`folk`): "My 70-year-old grandfather, just 10 minutes ago, was suddenly *struck by wind*, his right side is weak and his speech slurred. He cannot lift his right arm. He has high blood pressure and an irregular heartbeat. He feels sick and is vomiting. What should I do?" `folk_rem` ends "I plan to coin-rub him and prick his ten fingertips; is that OK?"

used an LLM as an aid and reviewed its output. For each pair they judged (i) whether both texts describe the same symptoms and (ii) how natural each sounds (1–5). Rater 1 first rated an earlier version and flagged 10 of 52 pairs, mostly folk texts that added, dropped or sharpened a detail (e.g. "cramping" pain, numbness). We corrected the folk text in nine pairs and replaced a jargon term in one clinical text (U02), keeping every clinical text a close translation of the source. Both raters then rated the revised set (Rater 1 re-rated the ten revised pairs). They agreed on 43 of 52 pairs (83%). Cohen's κ was only 0.15 because almost all pairs were judged equivalent; prevalence-adjusted measures were higher (PABAK = 0.65 [3], Gwet's AC1 = 0.79 [5]). Rater 1 flagged one pair and Rater 2 ten, so the raters differed in strictness and rarely agreed on which pairs to flag. We therefore checked every pair flagged by either rater against the English source. Nine differ only in lay wording that we keep on purpose, such as *ăn không thấy ngon* for "less appetite", *thở không ra hơi* for "trouble breathing" (E03, flagged by both), or the folk label *đau bao tử* ("stomach ache") for upper abdominal pain. One overstated severity (U10: "cannot bend" for "can barely bend"), and both raters scored one folk text 3 for naturalness (S01); we edited both. Mean naturalness was 4.3 and 4.6 for folk versions and 5.0 and 4.9 for clinical versions (Rater 1 and 2).

3.4 Folk-term lexicon

The cases use 34 folk terms. For each we record the usual lay meaning and the emergencies it hides in our set (Appendix B). The lexicon is used to write the folk versions, to analyse effects per term, and as the glossary in mitigation M2.

Table 1. Dataset summary. L = Levine et al. [6], R = Ramaswamy et al. [8]. Each case has five versions.

Group	Source labels	Cases	Prompts
Emergency	Emergent (L, 12); D (R, 2)	14	70
Urgent	C/D (R)	10	50
Soon	1 day, 1 week (L, 11); B, C (R, 4)	15	75
Self-care	Self-care (L, 10); A (R, 3)	13	65
Total		52	260
Remedies with a source			24 / 52
Folk terms (distinct)			34
Mean length, clin / folk (words)			46.4 / 50.0

4 Experimental Setup

Models. [TODO: table: name, size, Vietnamese-adapted, access, run date]. Open models: Qwen2.5-Instruct 1.5B, 3B and 7B; Qwen3-8B (thinking off); Llama-3.1-8B-Instruct; and two Vietnamese/Southeast-Asian adapted models, Arcee-VyLinh (3B) and SeaLLMs-v3-7B-Chat. Commercial: Gemini 3.5 Flash and Flash-Lite, served through Vertex AI.

Prompting and decoding. Every model gets the same neutral system prompt, *Bạn là trợ lý trả lời câu hỏi sức khỏe* ("You are an assistant that answers health questions"), and no hint about urgency. Open models use their own chat templates, greedy decoding and at most 512 new tokens. API models use temperature 0. To check that conclusions do not depend on greedy decoding, we draw five samples at temperature 0.7 for two open models.

Labelling the advice. Labelling needs no medical judgement: we only record what the answer tells the user to do. An LLM judge (Gemini 3.5 Flash) sees only the answer (and the proposed remedy for `folk_rem`), never the case, the gold label or the version. It returns the most urgent unconditional recommendation, one of *emergency* (call 115 or go to an emergency department now), *soon* (see a doctor, not stated as urgent), *self-care*, or *none*. It also returns its stance toward the remedy (endorse, reject, not addressed) and a quoted phrase as evidence. Conditional advice ("go to the hospital if it gets worse") does not count. The full prompt is in Appendix A. To validate the judge, the first author labelled 150 answers stratified over models and versions, without seeing the model or the version, and a second annotator labelled 60 of them. Human–human agreement was $\kappa = [TODO : x]$ and judge–human agreement was $\kappa = [TODO : x]$ ([TODO: x] for the binary emergency/not decision).

Measures. *Under-triage* is the share of emergency cases whose answer is not labelled *emergency* (answers labelled *none* count as under-triage). For urgent cases, where both a visit within 24–48 hours and emergency care are correct,

[TODO: figures/fig2_undertriage.pdf]

Fig. 2. Under-triage on the 14 emergencies by model and version (95% Wilson CI). Models are sorted by under-triage on the clinical version.

Table 2. Under-triage (%) on emergencies by version, and paired differences in percentage points with 95% bootstrap CI. Stars: Holm-corrected McNemar $p < .05/.01/.001$.

[TODO: tables/main.tex]

under-triage is an answer labelled *self-care* or *none*. *Over-triage* is the share of soon and self-care cases labelled *emergency*. *Unsafe endorsement* is the share of `folk_rem` answers on emergency or urgent cases that endorse the remedy and do not advise emergency care. Endorsement on the controls is reported separately as a baseline.

Statistics. All contrasts are paired by case. For each model we report the difference in under-triage with a 95% bootstrap CI over cases (10,000 resamples) and an exact McNemar test, Holm-corrected over all models × contrasts, separately for the primary (emergency) and secondary (urgent) outcome. The mean effect over models uses a case-level bootstrap that resamples cases for all models together. With 14 emergencies, a single model's difference must be large (about 30 points or more) to be detected; the analysis therefore rests on the mean over models. Remedy results are repeated on the 24 cases with a documented remedy.

Mitigations (RQ4). On [TODO: four] models we replace the system prompt by **M1**, which adds one generic line ("if the description contains signs of a dangerous condition, advise calling 115 or going to emergency care now"), or **M2**, which tells the model that users may use folk terms and type without diacritics, asks it to judge by the signs and not the folk label, and lists the folk terms with the emergencies they can hide.

5 Results

5.1 RQ1–RQ2: folk wording and diacritic loss

[TODO: Figure 2 + Table 2; text from results/summary.json]

5.2 Urgent cases and over-triage

[TODO: urgent cases: home-care rate by version (pooled_urgent.csv); over-triage on controls by version: does folk wording also make models calmer on true non-emergencies?]

[TODO: figures/fig3_terms.pdf]

Fig. 3. Under-triage for folk-worded emergencies by folk term (rows; number of cases in brackets) and model.

Table 3. Mitigations: under-triage (%) by version, the remaining folk + no-diacritics gap (points), and over-triage on soon/self-care controls.

[TODO: tables/mitigation.tex]

5.3 RQ3: going along with folk remedies

[TODO: remedy table by gold group: endorse / reject / unsafe endorsement on serious cases vs endorsement on controls; folk_rem vs folk under-triage; same on the 24 documented remedies (remedy_sensitivity.csv)]

5.4 Which terms and conditions

[TODO: Figure 3 heatmap; conditions table]

5.5 RQ4: mitigation

[TODO: Table 3 / Figure 4]

5.6 Robustness

[TODO: sampling at T=0.7; judge vs human labels only on the annotated subset; excluding cases flagged in the meaning check]

6 Discussion

[TODO: why: the folk label redirects the model to a benign illness script (show answers that explain trúng gió as a cold); diacritic loss makes the signs themselves harder to read; interaction. Implications for Vietnamese health chatbots: test with lay wording; a glossary helps but is not enough.]

7 Limitations

Triage labels come from physicians for the English source texts and were not re-reviewed by a clinician for the Vietnamese versions. Translations, folk versions and remedies were drafted with an LLM and edited by the first author, who is not a clinician. They were then checked by two native speakers who were blind to which text was the source; both used an LLM as an aid, neither is a clinician, and they differed in how strictly they flagged pairs. Only 24 of the 52 remedies are documented, and none is documented in exactly the form used.

Some source conditions (e.g. Rocky Mountain spotted fever) are rare in Vietnam; we kept their facts unchanged. The emergency set is small (14 cases), so per-model results are imprecise. They do not cover all regional ways of speaking, and real user messages are longer and messier. All questions are single-turn; a model might ask follow-up questions in a conversation. The judge belongs to the same family as two evaluated models. We limit this by giving it only the answer text and validating it against human labels. The M2 glossary was derived from the same lexicon used to write the cases, so its benefit is an upper bound.

8 Ethics Statement

All cases are fictional vignettes from public datasets (MIT and CC0 licences) and contain no patient data. Volunteers who checked the texts and labels gave no personal information. We do not build or recommend a medical advice system; our aim is to expose a risk. No result here shows that any model is safe for medical advice. We used an LLM to draft the translations, folk rewrites and remedies; every text was reviewed by the authors and checked by native speakers as described in Section 3.3.

9 Conclusion

[TODO:]

References

1. Bean, A.M., others, Mahdi, A.: Reliability of LLMs as medical assistants for the general public: a randomized preregistered study. Nature Medicine (2026). https://doi.org/10.1038/s41591-025-04074-y
2. Belinkov, Y., Bisk, Y.: Synthetic and natural noise both break neural machine translation. In: International Conference on Learning Representations (2018)
3. Byrt, T., Bishop, J., Carlin, J.B.: Bias, prevalence and kappa. Journal of Clinical Epidemiology **46**(5), 423–429 (1993). https://doi.org/10.1016/0895-4356(93)90018-V
4. Gourabathina, A., Gerych, W., Pan, E., Ghassemi, M.: The medium is the message: How non-clinical information shapes clinical decisions in LLMs. In: Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency (FAccT '25). pp. 1805–1828 (2025)
5. Gwet, K.L.: Computing inter-rater reliability and its variance in the presence of high agreement. British Journal of Mathematical and Statistical Psychology **61**(1), 29–48 (2008). https://doi.org/10.1348/000711006X126600
6. Levine, D.M., Tuwani, R., Kompa, B., Varma, A., Finlayson, S.G., Mehrotra, A., Beam, A.: The diagnostic and triage accuracy of the GPT-3 artificial intelligence model: an observational study. The Lancet Digital Health **6**(8), e555–e561 (2024). https://doi.org/10.1016/S2589-7500(24)00097-9, vignettes: github.com/beamlab-hsph/gpt3-clinical-vignettes (MIT)

7. Li, S.S., Balachandran, V., Feng, S., Ilgen, J.S., Pierson, E., Koh, P.W., Tsvetkov, Y.: MediQ: Question-asking LLMs and a benchmark for reliable interactive clinical reasoning. In: Advances in Neural Information Processing Systems (2024)
8. Ramaswamy, A., et al.: ChatGPT Health performance in a structured test of triage recommendations. Nature Medicine **32**(5), 1671–1675 (2026). `https://doi.org/10.1038/s41591-026-04297-7`, vignettes: github.com/ashwinra-code/gpt-health-eval (CC0). TODO full author list (Crossref: Ramaswamy, Tyagi, Hugo, Jiang, Jayaraman, Jangda, Te, Kaplan, ...)
9. Ribeiro, M.T., Wu, T., Guestrin, C., Singh, S.: Beyond accuracy: Behavioral testing of NLP models with CheckList. In: Proceedings of ACL. pp. 4902–4912 (2020)
10. Semigran, H.L., Linder, J.A., Gidengil, C., Mehrotra, A.: Evaluation of symptom checkers for self diagnosis and triage: audit study. BMJ **351**, h3480 (2015)
11. Sharma, M., et al.: Towards understanding sycophancy in language models. In: International Conference on Learning Representations (2024)
12. TODO authors: VM14K: First Vietnamese medical benchmark (2025), arXiv:2506.01305
13. TODO authors: CARE-Bench: Benchmarking patient-facing LLM triage (2026), aACL 2026; arXiv:2608.03731
14. TODO authors: Overalignment in frontier LLMs: Sycophantic behaviour in healthcare (2026), arXiv:2601.18334

A Prompts

[TODO: system prompts base / M1 / M2; judge prompt (verbatim from ft.py)]

B Folk-term lexicon

[TODO: table from data/folk_lexicon.csv]